

Project 2: N-Queens

Adapted from the [Course Scheduling of Stanford CS221](#)

Introduction

In this project, you are going to construct a CSP for N-Queens problem. You will be given several python files.

Files you'll edit:

[submission.py](#)

Where you need to fill in codes in three places.

Files you need to look at:

[test.py](#)

You can use this file to grade your program.

[run.py](#)

You can use this file to run your program to see results.

[csp.py](#)

The important class CSP is provided in this file.

Files to Edit and Submit: You will fill in portions of [submission.py](#) during the assignment. You should submit this file with your code and comments. Please do **NOT** change the other files in this distribution or submit any of our original files other than these files.

Evaluation: Your code will be auto-graded for technical correctness. Please do **NOT** change the names of any provided functions or classes within the code, or you will wreak havoc on the auto-grader. However, the correctness of your implementation -- not the auto-grader's judgements -- will be the final judge of your score. We will review and grade assignments individually to ensure that you receive correct credit for your work.

Academic Dishonesty: We will be checking your code against other submissions in the class for logical redundancy. If you copy someone else's code and submit it with minor changes, we will know. We trust you all to submit your own work only; *please* don't let us down. If you do, we will pursue the strongest consequences available to us.

CSP for N-Queens

A basic backtracking search [BacktrackingSearch](#) is already implemented in [submission.py](#). The function will help you solve the CSP. You should read it carefully to make sure that you understand how the backtracking search is working.

- a) [5 points] Let's create a CSP to solve the N-Queens problem: Given an $n \times n$ board, we'd like to place n queens on this board such that no 2 queens are on the same *row*, *column*, or *diagonal*. Implement [create_n_queens_csp\(\)](#) by **adding n variables** and some number of binary factors. Note that the solver collects some basic statistics on the performance of the algorithm. You should take advantage of these statistics for debugging and analysis. You should get 92 (optimal) assignments for $n = 8$ with exactly 2057 operations (number of calls to [backtrack\(\)](#)).

Hint: If you get a larger number of operations, make sure your CSP is minimal. Try to define the variables such that the size of domain is $O(n)$.

- b) [5 points] You might notice that our search algorithm explores quite a large number of states even for the 8×8 board. Let's see if we can do better. One heuristic we discussed in class is using most constrained variable (MCV): To choose an unassigned variable, pick the X_j that has the fewest number of values a which are consistent with the current partial assignment (a for which `check_factor()` on $X_j = a$ returns True). Implement this heuristic in `get_unassigned_variable()` under the condition `self.mcv = True`. It should take you exactly 1361 operations to find all optimal assignments for 8 queens CSP — that's 30% fewer!

Some useful fields:

- `csp.unary_factors[var][val]` gives the unary factor value.
- `csp.binary_factors[var1][var2][val1][val2]` gives the binary factor value.
 - ✧ Here, `var1` and `var2` are variables and `val1` and `val2` are their corresponding values.
- In `BacktrackingSearch`, if `var` has been assigned a value, you can retrieve it using `assignment[var]`. Otherwise `var` is not in `assignment`.

- c) [10 points] The previous heuristics looked only at the local effects of a variable or value. Let's now implement arc consistency (AC-3) that we discussed in lecture. After we set variable X_j to value a , we remove the values b of all neighboring variables X_k that could cause arc-inconsistencies. If X_k 's domain has changed, we use X_k 's domain to remove values from the domains of its neighboring variables. This is repeated until no domains have changed. Note that this may significantly reduce your branching factor, although at some cost. In `backtrack()` we've implemented code which copies and restores domains for you. Your job is to fill in `arc_consistency_check()`.

With AC-3 enabled, it should take you 769 operations only to find all optimal assignments to 8 queens CSP — That is almost 45% fewer even compared with MCV!

Take a deep breath! This part requires time and effort to implement — be patient.

Hint 1: documentation for `CSP.add_unary_factor()` and `CSP.add_binary_factor()` can be helpful.

Hint 2: although AC-3 works recursively, you may implement it iteratively. Using a queue might be a good idea. `li.pop(0)` removes and returns the first element for a python list `li`.

Submission (IMPORTANT THINGS! READ THREE TIMES)

- Submit your code, **only one file: `submission.py`**, through e-learning in **one** zip file with name “pj2-<student_id>.zip”. E.g. “pj2-15307130376.zip”
- **Project 2 is due on April 23rd, Tuesday, 11:55pm, 2019.**
- If you have any questions about this project, please contact the TA in charge Gao, Yixu via 15307130376@fudan.edu.cn.